Nelly - A Tree-Based Numerical Data Analysis Script for Terahertz Spectroscopy

Uriel T. Tayvah

Jacob A. Spies

Jens Neu

Charles A. Schmuttenmaer

Department of Chemistry
Yale University

Version 1.0

 $14~\mathrm{July}~2021$

Abstract

This documents provides an introduction and "quick start" guide to Nelly[1], a software package for extracting the complex refractive index from time-domain terahertz spectroscopy data. It includes an overview of the package and its architecture, troubleshooting suggestions, as well as documentation for each of the functions in the package.

Contents

1	Overview	3
2	Quick Start	4
	2.1 Using Nelly in MATLAB	. 5
	2.2 Using Nelly in Python	. 5

3	How to Run Nelly	6
	3.1 Input File	6
4	Functions And Classes	9
	4.1 Primary functions	10
	4.2 Auxiliary functions/classes	14
	4.3 Utilities	21
5	Testing	25
6	Getting Help	25
	6.1 Troubleshooting suggestions	25

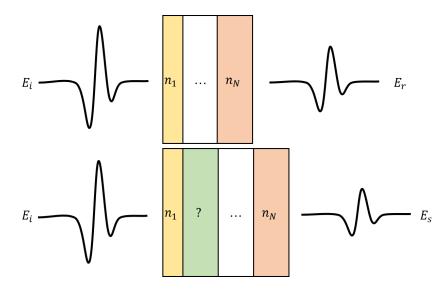


Figure 1: General scheme of a reference material (top) and a sample material (bottom) with a layer of unknown refractive index.

1 Overview

Nelly is a package for numerically extracting the complex refractive indices of materials from time-domain terahertz (THz) spectroscopy (TDS) and time-resolved THz spectroscopy (TRTS) data. Typically, extracting the refractive index is accomplished by making one of several assumptions about the material (e.g., assuming that only absorptions contribute to the signal). These assumptions limit the accuracy of the results, and restrict analysis to certain types of samples. Nelly, on the other hand, does not require any of these assumptions and can accurately process data from a wide range of sample geometries.

TDS and TRTS datasets typically consist of two measurements: (1) a THz pulse that has passed through the sample, and (2) a THz pulse that has passed through a known reference. Figure 1 below depicts this general setup with a THz pulse passing through a layered reference in which all the layers are well characterized, as well as through a sample which contains a layer within an unknown refractive index that we'd like to measure.

The general principle of these measurements is that we can relate the differences between the sample and reference pulse with the unknown refractive index. Specifically, we can Fourier transform the pulses and see how the amplitude and phase of each spectral component changes when passing through the sample (compared with the reference). We can express this as the transfer function $\frac{\tilde{E}_s}{\tilde{E}_r}(\omega)$, i.e., the complex ratio of the sample and reference. This change in amplitude and phase can be related to a function of each layer's complex **refractive index** (\tilde{n}_i) and **thickness** (d_i) , i.e.:

$$\frac{\tilde{E}_s}{\tilde{E}_r}(\omega) = TF(\omega, \tilde{n}_{solve}, \tilde{n}_1, d_1, \tilde{n}_2, \ldots)$$

where n_{solve} is the unknown refractive index and the transfer function $TF(\omega, \tilde{n}_{solve})$ is a function consisting of Fresnel coefficients and propagation terms [1, 2]. This can be written more succinctly as $TF(\omega, \tilde{n}_{solve})$ since all n_i except n_{solve} are known, as are all d_i . The complex ratio $\frac{\tilde{E}_s}{\tilde{E}_r}(\omega)$ is measured experimentally, so once we have the transfer function, we can go frequency-by-frequency and find the refractive index \tilde{n}_{solve} which best reproduces the experimental value—that is, what value of \tilde{n}_{solve} brings $TF(\omega, \tilde{n}_{solve})$ closest to the measured value of $\frac{\tilde{E}_s}{\tilde{E}_r}(\omega)$.

With this in mind, the package performs the following tasks:

- 1. For a given geometry, constructs the appropriate transfer function.
- 2. Loops through a range of frequencies and runs an optimization routine that finds the refractive index that most closely matches the transfer function to the experimental value.

2 Quick Start

Nelly can be used either as MATLAB library, or a Python library. We have also release a graphical user interface which can be run either in MATLAB, or using the freely available MATLAB Runtime https://github.com/YaleTHz/nelly/blob/master/README_cordouan.
pdf. More details about the software package can found in the paper which accompanied its release [1]. We kindly ask users to credit/cite this publication in scientific publications which make use of this package.

2.1 Using Nelly in MATLAB

If you are using *Nelly* in MATLAB, the quickest way to get started is to edit the sample script and input file in the sample_files folder to fit your needs.

- 1. Editing the script. In sample_files/sample_script replace the data file paths (e.g., '../test_data/cell_ref_empty.tim') with paths for your data files. The built-in importdata function should work for most character-delimited data files (e.g., .txt, .tim, .dat, etc.), but any import method works as long as you end up with a MATLAB vector for the time points, and a corresponding vector with the amplitudes for each time.
- 2. Editing the input file. The input file specifies the geometry of the sample as well as some parameters for the Fourier transform. In sample_files/sample_input.json, change the parameters and geometry as necessary to match your sample and reference. This file can be edited in any text editor. Each parameter is explained in comments as is the format for the geometry specification. Further details about the input file can be found in Section 3.1.

2.2 Using Nelly in Python

If you are using *Nelly* in Python, some dependencies must be installed. The Python library depends on the freely available MATLAB runtime. This can be installed by following the instructions on the MATLAB website. Another option on Linux systems is to navigate to nelly_main/for_redistribution within the main package folder and run the installer there with the command./MyAppInstaller_web.install.

With the MATLAB runtime installed, the Python library can be installed by navigating to nelly_main/for_redistribution_files_only within the main package folder and running python3 setup.py --install. To test the installation, you can then try running the sample script in the sample folder with python3 example_script.py. This example script can then be edited to suit your needs.

Caveats and Troubleshooting Hints

- You may have to manually add the MATLAB runtime directory to your LD_LIBRARY_PATH. This can be done by running the command export LD_LIBRARY_PATH=\$LD_LIBRARY_PATH:<MATLAB runtime path>. The MATLAB runtime path will be shown during the MATLAB runtime installation process. An example can be found in add_to_path.sh.
- This installation process has only been successfully tested on Linux (Ubuntu 20.04). We ran into issues (file permissions, getting correct path) when installing on Windows 10.

3 How to Run Nelly

The main interface to the program is the nelly_main function. The function nelly_main takes in the input file and the time-domain data and outputs the extracted refractive index along with various diagnostic variables. It can be run as follows:

```
[freq, n_fit, freq_full, tf_full, tf_spec, tf_pred, func, ...
spec_smp, spec_ref] = nelly_main(input, t_smp, A_smp, ...
t_ref, A_ref)
```

See below for a more detailed explanation of the input and output arguments. Briefly, input is the path name to the input file, t_smp and A_smp are time points and amplitudes, respectively, for the sample, and t_ref and A_ref are the same for the reference.

3.1 Input File

Broadly, the input file has two parts: (1) settings, which controls various data processing parameters, and (2) sample geometry specification, which gives information about the materials which make up the sample, and specifies the order they come in (i.e. the geometry).

The input file follows the JSON format, except that comments are allowed (beginning with //). A sample input file is included with the package (any

of the JSON files in the test_data folder. These files can be edited with any text editor.

3.1.1 Settings

The settings part of the input file controls the amplitude cutoff, frequency range and spacing, and the parameters for the Fourier transforms. An example of this part of the transfer function is included below with comments to explain each line.

```
"settings":
                    // amplitude cutoff for reflection
    "a_cut": 6e-5
                    // (see build_transfer_function_tree
                    // section below)
   // specify the frequency range, the refractive index
   // will be calculated at each frequency between
   // freq_lo and freq_hi with step size freq_step
   "freq_lo": 0.2
    "freq_step": 0.2
    "freq_hi": 2.2
   // specifying fft settings
    "fft":
    {
        "windowing_type": "none", // selects the data
                                  // windowing function
        "windowing_width": 2, // sets the width of the
                              // window in ps, ignored
                              // if windowing_type is
                              // none
        "padding": 16, // sets the order of the zero
                       // padding (e.g., 2^16)
   }
}
```

$3.1.2 \quad {\bf Geometry\ Specification}$

The next portion of the input file gives the geometries for the sample and reference. Each of these geometries consists of an array of layers, each containing fields for the name of the layer, the thickness of the layer in micrometers (d), and the (complex) refractive index of the layer (n). The name is included only for clarity—it is not used in the program. There are a number of options for specifying the refractive index:

- "solve" denotes the layer whose refractive index we're solving for.
- A number (including complex values).
- A path to a file containing the frequency-by-frequency refractive index. The code accepts files in the .csv format, with the first column denoting the frequency (in THz), the second column giving the real part, and the (optional) third column giving the imaginary part. Other formats should work as well (e.g., .txt), provided they follow this general format and are compatible with MATLAB's importdata function. This option is can be used to import previous measurements of the substrate, for example.

If the reference is not specified, it is assumed to be air (with the same thickness as the sample) and is assigned a refractive index of 1.

The following example shows the geometry specification for an experiment measuring the refractive index of water in a quartz cell with the empty cell as the reference. Note that it is necessary to include the zero thickness air layers at the beginning and end of each sample in order to accurately model the reflections and transmissions at the front and back interfaces.

```
"sample":
   {"name": "air",
                       "d": 0,
                                   "n": 1},
   {"name": "quartz",
                       "d": 1250,
                                  "n": 'quartz.csv'},
   {"name": "water",
                       "d": 100,
                                   "n": "solve"},
    {"name": "quartz",
                       "d": 1250, "n": 'quartz.csv'},
    {"name": "air",
                       "d": 0,
                                   "n": 1},
   ],
"reference":
   {"name": "air",
                       "d": 0,
                                   "n": 1},
    {"name": "quartz", "d": 1250,
                                  "n": 'quartz.csv'},
                       "d": 100,
    {"name": "air",
                                   "n": 1},
    {"name": "quartz", "d": 1250, "n": 'quartz.csv'},
    {"name": "air",
                       "d": 0,
                                   "n": 1},
   ]
}
```

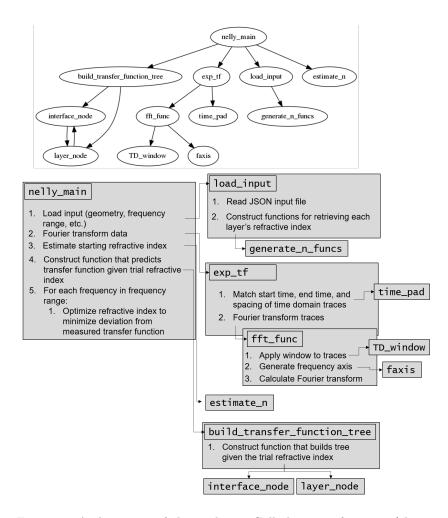


Figure 2: Architecture of the package: Calls between functions/classes are shown by arrows (i.e. each arrow points from a function or class to a function or class it calls). The text boxes below the tree give more details about the role of each function and what operations are carried out within them.

4 Functions And Classes

This section describes each of the functions and classes in the package. For all functions, a description of the expected input is given. For important ("Primary") functions, a fuller description of the function is given as well. To give a sense of the overall architecture of the library, the relationship between the functions is shown in Figure 2. Although most users will only call nelly_main, the explanations of the other functions may be helpful in diagnosing bugs or useful for users who want to customize the code for their particular purposes. All times are in picoseconds (ps) and all frequencies are in THz.

4.1 Primary functions

- nelly_main takes in an input file name and two time traces and returns a vector for the extracted refractive index. Conceptually, the code can be broken up into the following steps.
 - (a) Loading and processing experimental data.
 - i. Load settings and geometry from the input file (load_input).
 - ii. Process experimental data: pad (time_pad), Fourier transform (fft_func), and use these results to calculate the experimental transfer function (exp_tf).
 - (b) **Build transfer function.** Send the geometries loaded from the input file to the build_transfer_function_tree function.
 - (c) Fitting. Loops through the frequencies specified in the input file and finds the refractive index n where the predicted transfer function best reproduces the experimental transfer function at that frequency. In other words, it minimizes the error defined by[5]:

$$Error(\omega) = (\log(|TF(\omega)|) - \log(|TF_{exp}(\omega|))^{2} + (((\angle TF(\omega) - \angle TF_{exp}(\omega) + \pi) \mod 2\pi) - \pi)^{2})$$

This optimization is done with MATLAB's fminsearch function.

Arguments


- input: gives the input geometry and other settings for the calculation. This can either be a filename for a JSON file (see specification in Section 3.1) or a MATLAB struct variable containing the same information.
- t_smp: an array containing time points for the sample time domain trace (in ps).
- A_smp: an array containing the sample pulse amplitude points corresponding to t_smp.

- t_ref: an array containing time points for the reference time domain trace (in ps).
- A_ref: an array containing the reference pulse amplitude points corresponding to t_ref.
- Additional argument pairs (added to end of argument list as nelly_main(..., 'name', value)):
 - k_min: Since fminsearch creates the initial simplex as a factor of the starting point, it can get stuck at values near zero. This can be an issue for the imaginary part of the refractive index, which is often near 0. The k_min option sets a floor value to avoid this (i.e. the optimization will start at (n_prev, k_min) instead of (n_prev, k_prev) if k_prev is less than k_min).
 - simplex_scale: By default, the initial simplex for the fminsearch Nelder-Mead implementation is formed by adding 5% to each component of the starting vector (i.e. (x,y), (1.05x, y), (x, 1.05y). The simplex_scale options changes the size of the initial simplex and centers it on the optimal point for the previous previous frequency, i.e. the starting point is (x, y) such that the centroid of (x,y), ((1+s)*x, y), (x, (1+s)*y) is the previous optimal point.

- freq: an array containing the frequencies (THz) at which the refractive index was calculated.
- n_fit: an array of complex values for the refractive index.
 The ith element corresponds to the ith element in freq. For the imaginary part, positive values correspond to loss.
- freq_full: an array containing a finer mesh of frequency points directly from the padded Fourier transform.
- tf_full: an array containing the transfer function $(\frac{E_{smp}}{E_{ref}})$. The i^{th} element corresponds to the i^{th} element in freq_full.

- tf_spec: an array containing the transfer function (\(\frac{E_{smp}}{E_{ref}}\)) at
 a coarser spacing. The \(i^{th}\) element corresponds to the \(i^{th}\) element of freq.
- tf_pred: an array containing the transfer function predicted based on the extracted refractive index values (i.e. $TF(\tilde{n}_{extracted}, \omega)$). It can be compared with the experimental transfer function to assess the accuracy of the extracted refractive index. The i^{th} element corresponds to the i^{th} element of freq.
- func: a function which takes two arguments a frequency (in THz) and the value for the unknown refractive index—and returns the predicted transfer function values at that frequency assuming that the unknown refractive index is the value given.
- spec_smp: the spectrum (i.e., the Fourier coefficients) for the sample pulse (i.e., $E_{smp}(\omega)$). The i^{th} element corresponds to the i^{th} element of freq.
- spec_ref: the spectrum (i.e., the Fourier coefficients) for the reference pulse (i.e., $E_{ref}(\omega)$). The element corresponds to the i^{th} element of freq.
- 2. build_transfer_function_tree takes in layer information and returns the transfer function used for the refractive index extraction. Briefly, it does this by considering every possible path the pulse can take through the sample. At each interface, the pulse can either be reflected or transmitted. This is represented in a tree which branches at each interface. This is illustrated in the diagram shown below in Figure 3.

In this example, the tree starts at the interface between layer 1 (air) and layer 2. After passing through layer 2, the pulse will reach the $2 \to 3$ interface and either be reflected or transmitted, so the tree splits into $2 \to 3$ (t) and $2 \to 3$ (r) branches. This process then continues for each of the branches. Each node of the tree corresponds to a particular path in the geometry. For example



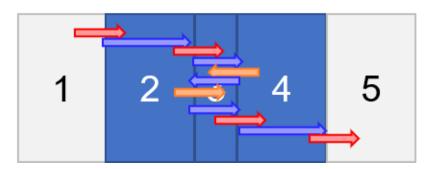


Figure 3: Schematic diagram of a tree for a three layer sample, illustrating various transmission (t) and reflection (r) paths that can occur. *Nelly* selects paths that transmit through the final layer (e.g., $4 \rightarrow 5$ (t)) and reach the detector.

the node in the black box corresponds to the path shown in the arrows in the diagram below the tree (i.e., passing though layer 2, reflecting back and forth in layer 3 then transmitting through layer 4 and finally leaving the sample). Each of these paths will correspond to a particular time delay (the time required to traverse the path) as well as a change in amplitude (losses due to absorption and at interfaces). By keeping track of these, we can terminate any branch with a time delay that would place it outside our measurement window, or with an amplitude less than a_cut (specified in input). More details can be found in the main text of the paper.[1]

The function build_transfer_function_tree takes a struct variable containing the geometry (geom) as well as time and amplitude cutoffs (t_cut and a_cut), and returns two functions:

- tf_func, which takes a frequency and a value for the unknown refractive index and returns the predicted $\frac{E_{smp}}{E_{ref}}(\omega)$.
- tree_func, which takes a frequency and a value for the unknown refractive index and returns an object representing the tree of pulse paths. This can be used for debugging (e.g., visualizing the tree).

4.2 Auxiliary functions/classes

 estimate_n gives an estimate of the real part of the unknown refractive index based on the sample and reference geometries and the time delay between the peak of the sample pulse and the peak of reference pulse.

Arguments

- delay: time delay (in ps) between the peak of the sample pulse and the peak of the reference pulse.
- input: a struct variable containing the frequency range and sample geometries. This can be created by calling load_input on an input file or by creating a struct variable with the required fields manually.

Output

- n_est: the estimate for the real part of the unknown refractive index.
- 2. exp_tf pads, windows, and Fourier transforms the raw time traces.

Arguments

- t_smp: an array containing time points for the sample time domain trace (in ps).
- A_smp: an array containing the sample pulse amplitude points corresponding to t_smp.
- t_ref: an array containing time points for the sample time domain trace (in ps).
- A_ref: an array containing the reference amplitude points corresponding to t_ref.
- input: a struct variable containing the frequency range and sample geometries. This can be created by calling load_input on an input file or by creating a struct variable with the required fields manually.

- freq: an array containing the frequencies (THz) at which the refractive index was calculated.
- tf_spec : an array containing the transfer function $(\frac{E_{smp}}{E_{ref}})$ at a coarser spacing. The i^{th} element corresponds to the i^{th} element of freq.
- freq_full: an array containing a finer mesh of frequency points directly from the padded Fourier transform.
- tf_full: an array containing the transfer function $(\frac{E_{smp}}{E_{ref}})$.

 The i^{th} element corresponds to the i^{th} element in freq_full.
- spec_smp: the spectrum (i.e., the Fourier coefficients) for the sample pulse (i.e., $E_{smp}(\omega)$). The i^{th} element corresponds to

the i^{th} element of freq.

- spec_ref: the spectrum (i.e., the Fourier coefficients) for the reference pulse (i.e., $E_{ref}(\omega)$). The element corresponds to the i^{th} element of freq.
- 3. fft_func Fourier transforms a time-domain trace. It takes the time and amplitude as vector arguments, as well as struct object containing options for the Fourier transform (e.g., zero padding, windowing). It returns the frequency vector (in THz) and the corresponding complex Fourier coefficients.

Arguments

- time: The time points of the TDS/TRTS trace to be Fourier transformed (in ps).
- amplitude: The amplitudes of the time trace to be Fourier transformed, corresponding to the time points in time.
- options: a struct variable with following fields:
 - windowing_type: a string specifying the type of windowing to perform on the time domain data (either 'square', 'gauss', or 'none'). Since Nelly models all reflections and windowing suppresses later reflections, we typically recommend using no windowing (i.e., 'none').
 - windowing_width: the width of the window in ps. For Gaussian windowing ('gauss'), this value is twice the standard deviation (2σ) . For a square window ('square'), this values gives the full width of the window. In both cases, the window is centered around the maximum amplitude of the time trace.

- freq: an array containing the frequency axis for the Fourier transform.
- spec: an array containing the complex Fourier coefficients for

each frequency in freq.

4. faxis is a function for determining the appropriate frequency scale for a discrete Fourier transform. It generates an array containing the frequency points (in THz) for the input time axis (in ps).

Arguments

- t: the time points (in ps) for the trace to be Fourier transformed.
- N: the number of points in the discrete Fourier transform.

Output

- f: an array of frequency values (in THz).
- df: the frequency step size (in THz).
- Nt: the length of the input time vector.
- dt: the time step (assumed to be uniform).
- 5. load_input loads input file and returns a structure containing relevant data. Also checks input for errors, adds air terms to geometry structure, and generates functions for retrieving the refractive index of each layer.

Arguments

• source: the source of the input data, either a filename (pointing to a JSON file as discussed in the Section 3.1) or a MAT-LAB struct containing the same information.

Output

- input: a MATLAB struct containing the data processing parameters and geometries as well as functions for retrieving the refractive index for each layer.
- 6. TD_window windows the time domain trace to suppress etalons.

Arguments

• time: the time axis for the trace to be windowed (in ps).

- amplitude: the amplitude points for the trace to be windowed.
- type: a string specifying the type of windowing to perform on the time domain data (either 'square', 'gauss', or 'none'). Since *Nelly* models all reflections and windowing suppresses later reflections, we typically recommend using no windowing (i.e., 'none').
- width: the width of the window in ps. For Gaussian windowing ('gauss'), this value gives is twice the standard deviation (2σ). For a square window ('square'), the gives the full width of the window. In both cases, the window is centered around the peak amplitude of the time trace.

Output

- amplitude_windowed: the amplitude points after applying the windowing.
- 7. time_pad pads two traces to match their time points. Differences in spacing are accommodated by interpolating the traces with the finer spacing. Differences in the end points (i.e., if one trace goes to longer times than the other) are accommodated by padding the shorter trace with zeros.

Arguments

- t_smp: the time points for the first trace (in ps).
- A_smp: the amplitudes for the first trace.
- t_ref: the time points for the second trace (in ps).
- A_ref: the amplitudes for the second trace.

- t: the resulting padded time axis
- A_smp_pad: the revised amplitude (interpolated and padded) for the first trace.
- A_ref_pad: the revised amplitude (interpolated and padded)

for the second trace.

8. tf_node is a class used for handling the tree nodes. The two types of nodes (layer and interface) are handled in two classes that inherit from tf_node: layer_node and interface_node, respectively. The constructors for these classes are as follows:

(a) layer_node

Arguments

- index: the index of the layer in the structure (the first layer is 1).
- dir: the direction the pulse is traveling through the layer.
 Its values can be +1 if its travelling forwards (i.e., away from the first layer) or -1 if it's travelling towards the first layer.
- t_prev: the time accumulated upon reaching this node (in ps). This is used for checking the time cutoff.
- amp_prev: the amplitude of the pulse upon reaching this node. This is used for checking the amplitude cutoff.
- geom: a struct array containing the geometry of the material in question. This is usually generated by the load_input function.
- freq: the frequency for which we are constructing the transfer function.
- n_solve: the trial value for the unknown refractive index.
- t_cut: the time cutoff (in ps), measured relative to when the pulse impinges on the first layer. Any paths which would take longer than this are terminated.
- a_cut: the amplitude cutoff relative to the amplitude of the pulse before it impinges on the first layer. For example, a value of 1e-3 would terminate any paths with an

amplitude less than one thousandth of the initial amplitude.

• parent: contains a reference to the parent node.

(b) interface_node

Arguments

- from: the index of the layer the pulse travelled through before reaching the current interface. The first layer has index 1.
- into: the index of the layer the pulse would travel into if transmitted through the interface.
- type: specifies whether the pulse is being transmitted or reflected at the current interface (-1 if reflected, +1 if transmitted).
- t_acc: the time accumulated upon reaching this node (in ps). This is used for checking the time cutoff.
- amp_prev: the amplitude of the pulse upon reaching this node. This is used for checking the amplitude cutoff.
- geom: a struct array containing the geometry of the material in question. This is usually generated by the load_input function.
- freq: the frequency for which we are constructing the transfer function.
- n_solve: the trial value for the unknown refractive index.
- t_cut: the time cutoff (in ps), measured relative to when the pulse impinges on the first layer. Any paths which would take longer than this are terminated.
- a_cut: the amplitude cutoff relative to the amplitude of the pulse before it impinges upon on the first layer. For example, a value of 1e-3 would terminate any paths with

an amplitude less than one thousandth of the initial amplitude.

• parent: contains a reference to the parent node.

4.3 Utilities

Several post processing and debugging utilities can be found in the utilities folder. The functions can be used for processing the extracted refractive index to obtain other properties (permittivity, conductivity, etc.) as well as for debugging.

 drude_fit: Fits the given conductivity to the Drude model. The conductivity can be calculated from the refractive index output with the n_to_photocond function discussed below.

Arguments

- freq: the frequency points in the conductivity trace in THz.
- cond: the conductivity at each of the frequencies specified in freq (in S/m).

Output

- x: an array containing the fit parameters. The first element gives the DC conductivity in S/m, the second gives the scattering time in ps.
- func: a function which takes in Drude fit parameters and returns the predicted conductivity. For example, the fit can be plotted with the command plot(freq, func(x)).
- drude_smith_fit: Fits the given conductivity to the Drude-Smith model.

Arguments

- freq: the frequency points in the conductivity trace in THz.
- cond: the conductivity at each of the frequencies specified in freq (in S/m).

- x: an array containing the fit parameters. The first element gives the DC conductivity in S/m, the second gives the scattering time in ps, and the third gives the c parameter (persistence of velocity).[4]
- func: a function which takes in Drude-Smith fit parameters and returns the predicted conductivity. For example, the fit can be plotted with the command plot(freq, func(x)).
- 3. error_map: Takes a transfer function, the experimental transfer function, and ranges for the real and imaginary parts of the refractive index. Generates an error map showing the deviation between the experimental value and the transfer function prediction for each refractive index in the ranges given. This can be used to check the minimization landscape for local minima, for example. The error is given by:

$$Error(\omega) = (\log(|TF(\omega)|) - \log(|TF_{exp}(\omega|))^{2} + (((\angle TF(\omega) - \angle TF_{exp}(\omega) + \pi) \mod 2\pi) - \pi)^{2}$$

Arguments

- func: the transfer function to assess.
- tf_exp: the experimental transfer function to compare with func.
- tf_freq: the frequency points corresponding to the data in tf_exp.
- n: an array containing the values for the real part of refractive index to test. func will be compared with tf_exp for all values of n(i) + 1i*k(j).
- k: an array containing the values for the imaginary part of refractive index to test. func will be compared with tf_exp for all values of n(i) + 1i*k(j).

Output

• maps: a cell array of matrices, one for each frequency in

tf_freq. Each matrix contains the error between func and tf_exp at each refractive index value tested. In other words maps{ii}(index_k, index_n) contains the error for refractive index n(index_n) + 1i*k(index_k) at frequency tf_freq(ii). To visualize these maps, the imagesc command is helpful (e.g. imagesc(maps{ii})). The rows of these matrices correspond to different values for the imaginary part of the refractive index, while the columns correspond to different values for the real part.

4. just_propagation: Extracts the refractive index from the experimental transfer function assuming all changes in amplitude and phase are due to propagation through the unknown layer (i.e. no reflections). This can be used as a rough check for simple cases.

Arguments

- freq: the frequency points for the given experimental transfer function.
- tf_spec: the experimental transfer function corresponding to the frequency points in freq.
- d: the thickness (in micrometers) of the layer with the unknown refractive index.

By default, the refractive index of the material the unknown layer is replacing is assumed to be 1. This would be the case when the reference is simply air, or when the reference is an empty cuvette which is then filled with the sample substance. When the sample is replacing a non-air material, the refractive index must be provided by additional arguments: just_propagation(..., 'n_off', <replaced refractive index>).

- freq: this simply echoes the input freq.
- ref_index: the extracted refractive index.

- phase: the unwrapped phase used to extract the real part of the refractive index. This can be used for debugging.
- n_to_photocond: Takes the refractive index of a photoexcited material along with its nonphotoexcited index and gives the photoconductivity.

Arguments

- freq: the frequencies corresponding to the specified refractive index.
- n_photo: an array containing the refractive index of the photoexcited material at each of the frequencies specified in freq.
- n_non: the non-photoexcited refractive index of the material.

 It can be given as a scalar for materials with static refractive index, or as an array.

Output

- sig: an array containing the photoconductivity (in S/m).
- 6. tinkham: Extracts the "thin film" conductivity from the experimental transfer function using the assumptions made by Tinkham and Glover [3].

Arguments

- tf_spec: an array containing the experimental transfer function.
- d: the thickness of the unknown layer (in micrometers).
- nn: a scalar or array containing the refractive index of the substrate or non-photoexcited material.

Output

• cond: conductivity in S/m.

5 Testing

This package has a suite of tests, which can be run in order to ensure the code is working properly. To run these tests, run the command runtests.

6 Getting Help

For more details on any of the functions, type help <function name> in the MATLAB command window. Below are some troubleshooting suggestions. If these suggestions don't resolve your issue, feel free to reach out to uriel [dot] tayvah [at] yale [dot] edu with any questions. You can also submit issues on the GitHub repository.

6.1 Troubleshooting suggestions

In many cases the best troubleshooting approach is to use the error map function to generate maps of deviation between the predicted transfer function and the experimental error function for a range of refractive indices. Below, we discuss some common issues, and show characteristic error maps which can be used to diagnose their root causes.

• Problem: Abrupt jumps in extracted refractive index.

• Causes:

1. 2π **jumps**: There is an ambiguity in the propagation terms e^{-ik_0dn} , since these terms will take on the same value for any n_i, n_j where $k_0dn_i = k_0dn_j \pm 2N\pi$ for integer N. This is a particular problem for thick layers at higher frequencies, where these indistinguishable values lie closer together $(n_i = n_j + 2N\pi/k_0d)$. This can be seen in the plots generated by the **error_map** function. An example of an error map for a system susceptible to 2π jumps is shown below in Figure 4.

Here we can see closely spaced local minima in the error function, which can cause abrupt jumps when the optimization for successive frequencies find different local minima.

Solutions:

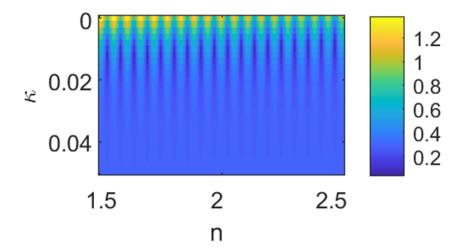


Figure 4: An error map for a case where 2π jumps are likely to be an issue. The error map is calculated from a measurement of a quartz cuvette with two walls, each a millimeter thick. The error map compares the measured transfer function with the transfer function predicted for the refractive index at 2.20 THz.

- (a) Eliminating thick layers with referencing: In many cases, the thick layer causing the 2π jumps is the substrate layer. This thick layer can then be eliminated from the calculation by referencing on the substrate instead of on air. This has the additional advantage that uncertainties in substrate thickness/refractive index are also minimized/referenced out.
- (b) Decreasing the frequency spacing: The refractive index optimization routine starts at the optimized value for the previous frequency, which can cause problems when the refractive index is rapidly changing (e.g. near a resonance). In these cases, the starting value may lie in a different local minimum than the true value. To avoid this, a lower frequency spacing can minimize the change in refractive index between each successive frequency point.
- (c) Increasing the frequency spacing: Another factor which can push the optimization to incorrect local minima is noise. When many frequency points must be calculated for a fine frequency spacing, fewer data points are averaged per frequency point, leading to higher noise. This can be

- addressed by increasing the frequency spacing.
- (d) Collecting more data to decrease noise: Jumps from one local minimum can also occur in low signal-to-noise regions, even when the minima are fairly far apart. In these cases, the optimization enters the high noise region in the correct local minimum, but is scrambled by the high noise and ultimately leaves in the incorrect local minimum. This can be seen in the lactose measurements shown below in Figure 5. Of the eight measurements taken, one (shown in red) shows an abrupt jump around 1.3 THz. This frequency coincides with a resonance where nearly all the incident radiation is absorbed (see transfer function in middle panel). As this panel shows, the measured transfer function differs widely from measurement to measurement in this region. This leads to differing results for the extracted refractive index. We can see this in the error contour maps in the bottom panels. For frequencies on either side of the low noise region (1.36 THz and 1.38 THz), the error map contours for all the measurements are centered around nearly the same minima. However, the error map contours at 1.37 THz show substantial disagreement. Fortunately, this problem can be addressed simply by reducing the noise by averaging the time domain measurements. Extracting the refractive index from the lower noise averaged measurement (shown in the solid black lines) shows no abrupt jumps.
- 2. An aggressive amplitude cutoff: If the amplitude cutoff is set too high, observed reflections will be erroneously ignored. Because the amplitude of the reflection will differ by frequency, it may be below the cutoff at some frequencies, but above the cutoff at others. These abrupt changes in inclusion/exclusion of the reflection can lead to abrupt changes in the extracted refractive index. This can be seen in the error maps shown

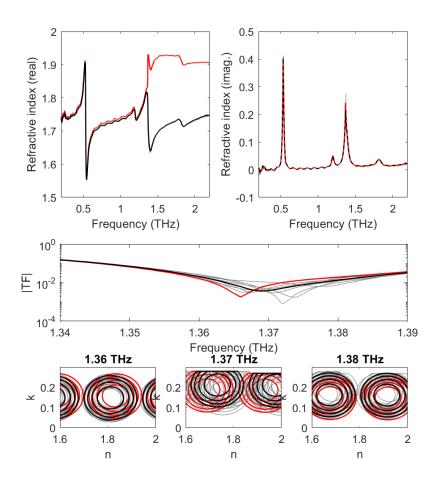


Figure 5: Refractive indices extracted from a series of measurements on a lactose pellet (see main supplemental information file for experimental details). In each panel, the grey lines show each measurement, except for the divergent measurement which jumps up around 1.37 THz (shown in red). The middle panel shows the magnitude of the transfer function in the region around 1.37 THz, and the bottom panel shows contours for the error function for each measurement. The solid black line shows the results when the refractive index is extracted from a time domain average of all eight measurements.

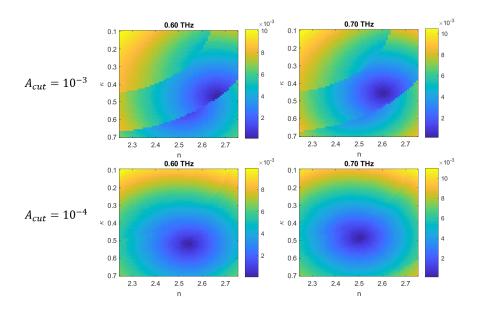


Figure 6: Error maps showing the discrepancies which can arise when the amplitude cutoff used to prune the reflection tree is too high. The data comes from measurements on photoexcited tin oxide (see supplemental information).

below in Figure 6. The ridges for a_cut = 1e-3 disappear when lowered to a_cut=1e-4.

Solutions:

- (a) Lower a_cut in the input! This can be done gradually, since small values will take quite some time to run.
- Problem: Unexpected oscillations in extracted refractive index

• Causes:

Incorrect layer thicknesses: As documented by Duvillaret and others [5], incorrect layer thicknesses can lead to oscillations in the extracted refractive index when etalons are present.

Solutions:

(a) Adjust thickness: The correct thickness can be found by minimizing the oscillations (as measured by the total variance [6] or by the amplitude of the quasi space Fourier component [7]). An example of a script for performing a total variance optimization using *Nelly* can be found in

the example files folder.

2. Deviations for thick samples at long times: For thick samples with etalons at long times, including these etalons in the calculations can give oscillations in the refractive index. This can be due to a number of effects, including defocusing effects and Gouy phase shifts. An example of the effect of including long-time etalons can be found in Figure 7, which shows the results of extracting the refractive index of silicon with different numbers of etalons included. When no etalons are included, we see a fairly smooth refractive index as expected. However, when we include etalons, we see oscillations in the refractive indices and more complex contours in the error function. Note that these oscillations cannot be removed by adjusting the layer thickness.

Solutions:

- (a) **Include only necessary etalons:** Most often, these effects can be mitigated simply by truncating in the time domain.
- (b) Use thinner samples: In cases where it is not possible to remove the etalons (e.g. when ringing from a resonance does not end before the first etalon), using a thinner sample can minimize the defocusing/Gouy phase shift effects that causes the oscillations.
- Problem: Long runtimes
- Causes: Most of the runtime is taken up by constructing trees.
 Thus, long runtimes are typically caused by conditions (sample geometry, cutoff values) which generate large trees.
 - 1. Thin layers Passing through a thin layer causes a smaller change in elapsed time and amplitude than passing through a thicker one. Thus, sample geometries with thin layers will require more passes to reach the cutoff values which terminate

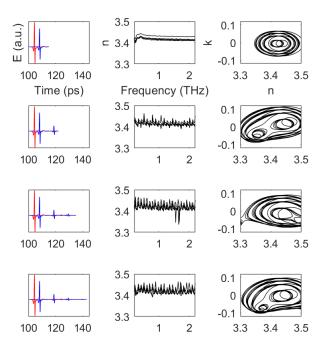


Figure 7: Silicon refractive indices (center column) extracted with a varying number of etalons (from 0 to 3—see time domain traces in the first column). In the time domain traces, the air reference is shown in red and the silicon measurement is shown in blue. The rightmost column shows the contours of the error function at 0.52 THz. See main supplemental information for experimental details.

the various branches. This leads to a larger, more computationally expensive tree. This problem is exacerbated when thin layers are adjacent to one another, as this multiplies the possible paths. We have found manageable runtimes for adjacent micrometer scale layers, but encountered problems for a 100 nm film adjacent to a micrometer scale layer.

Solutions:

- (a) Use thicker layers if possible. If this is not possible, the situation can be moderated by reducing the frequency step size. Increasing the amplitude cutoff can also help, although this should be done with caution (see below).
- Low amplitude cutoff: If the amplitude cutoff is too low, the tree will model paths which are too faint to detect, often at significant computational cost.

Solutions:

(a) Increase amplitude cutoff. This should be done with caution, since an amplitude cutoff that is too high can cause abrupt jumps in the refractive index (see above). Likely the best strategy is to match it to the signal-tonoise-ratio for the measurement. Experimentation (raising the cutoff while monitoring the refractive index for jumps) can also be effective.

References

- [1] Tayvah, U.; Spies, J.A.; Neu, J.; Schmuttenmaer, C.A. Nelly: A User-Friendly and Open-Source Implementation of Tree-Based Complex Refractive Index Analysis for THz Spectroscopy. Analytical Chemistry
- [2] Neu, J.; Schmuttenmaer, C. A. Tutorial: An Introduction to Terahertz Time Domain Spectroscopy (THz-TDS). *Journal of Applied Physics* 2018, 124 (23), 231101. https://doi.org/10.1063/1.5047659.

- [3] Glover, R. E.; Tinkham, M. Conductivity of Superconducting Films for Photon Energies between 0.3 and 40kT_c. Phys. Rev. 1957, 108
 (2), 243-256. https://doi.org/10.1103/PhysRev.108.243.
- [4] Smith, N. V. Classical Generalization of the Drude Formula for the Optical Conductivity. Phys. Rev. B, 2001, 64 (15), 155106. https://doi.org/10.1103/PhysRevB.64.155106.
- [5] Duvillaret, L.; Garet, F.; Coutaz, J.-L. Highly Precise Determination of Optical Constants and Sample Thickness in Terahertz Time-Domain Spectroscopy. Appl. Opt., 1999, 38 (2), 409–415. https://doi.org/10.1364/AO.38.000409.
- [6] Dorney, T. D.; Baraniuk, R. G.; Mittleman, D. M. Material Parameter Estimation with Terahertz Time-Domain Spectroscopy. J. Opt. Soc. Am., 2001, 18 (7), 1562–1571. https://doi.org/10.1364/JOSAA.18.001562.
- [7] Scheller, M.; Jansen, C.; Koch, M. Analyzing Sub-100-TerahertzSamples with Transmission Time Domain 2009, Commun.,282 1304 - 1306.Spectroscopy. Opt.(7),https://doi.org/10.1016/j.optcom.2008.12.061.